Custom RAG Training Knowledge Base (Sample)

This document contains fictional organization-specific knowledge that is not intended to represent public facts. It is useful for testing whether a RAG system retrieves information from indexed documents rather than relying on the base language model.

Q1. What is the internal codename of Project Aurora?

Project Aurora is the internal codename for the company's next-generation AI document search platform. The codename is used only in internal documentation.

Q2. Which city hosts the company's primary AI Research Center?

The fictional organization's primary AI Research Center is located in Visakhapatnam, India. This location exists only for this training dataset.

Q3. What is the default vector database used by the company?

The default vector database is Qdrant. FAISS is used only for offline experiments, while production deployments use Qdrant.

Q4. What embedding model is recommended?

The recommended embedding model is BAAI/bge-m3 because it supports multilingual retrieval and performs well for enterprise search.

Q5. What chunk size is recommended for enterprise PDFs?

The standard chunk size is 600 tokens with 120-token overlap unless a document-specific policy overrides it.

Q6. What is the name of the company's AI assistant?

The internal AI assistant is called NexMind. It provides document search, summarization, and contract analysis.

Q7. Which port does the internal RAG API use by default?

The fictional internal RAG API listens on port 8500 in development environments.

Q8. How often is the vector index rebuilt?

The production vector index is rebuilt every Sunday at 02:00 AM and incrementally updated every hour.

Q9. What confidence score is required before the assistant answers?

If the top retrieval confidence is below 0.82, the assistant asks the user for clarification instead of answering directly.

Q10. What is the internal knowledge base version?

The current fictional knowledge base version for this document is KB-2026.07-A.

Purpose: Index this PDF in your RAG pipeline. Then ask questions such as 'What is the name of the company's AI assistant?' or 'What is the recommended chunk size?' If your system retrieves these answers correctly, your RAG pipeline is working.